

CS 234 Winter 2019

Default Final Project

Estimation of the Warfarin Dose

Due date: Wednesday March 20 at 11:59pm PST.

Late days: You cannot use late day on the final project report and poster submission. Please read the [late day policy](#) on website for more details.

This project can be done in groups up to 3. We encourage you to do groups in 3, work together and discuss so that all group members have full understanding of the submitted work. You may use any existing code, libraries, etc. and consult any papers, books, online references, etc. for this project. However, you must cite your sources in your writeup and clearly indicate which parts of the project are your contributions and which parts were implemented by others. Under no circumstances may you look at another group's code or incorporate their code into your project.

Read the section on [Academic Collaboration and Misconduct](#) for an overview of the collaboration policy and academic integrity standards expected in general.

Contents

Introduction	2
Dataset	3
Baselines [10pts]	3
Implementing a Linear Bandit Algorithm on Warfarin Dataset [50pts]	4
Further Directions [40pts]	4

Introduction

Warfarin is the most widely used oral blood anticoagulant agent worldwide; with more than 30 million prescriptions for this drug in the United States in 2004. The appropriate dose of warfarin is difficult to establish because it can vary substantially among patients, and the consequences of taking an incorrect dose can be severe. If a patient receives a dosage that is too low, they may experience excessive anti-coagulation (which can lead to dangerous bleeding), and if a patient receives a dosage which is too high, they may experience inadequate anti-coagulation (which can mean that it is not helping to prevent blood clots). Because incorrect doses contribute to a high rate of adverse effects, there is interest in developing improved strategies for determining the appropriate dose ([Consortium, 2009](#)).

Commonly used approaches to prescribe the initial warfarin dosage are the *pharmacogenetic algorithm* developed by the IWPC (International Warfarin Pharmacogenetics Consortium), the *clinical algorithm* and a *fixed-dose* approach.

In practice a patient is typically prescribed an initial dose, the doctor then monitors how the patient responds to the dosage, and then adjusts the patient's dosage. This interaction can proceed for several rounds before the best dosage is identified. However, it is best if the correct dosage can be initially prescribed.

This project is motivated by the challenge of Warfarin dosing, and considers a simplification of this important problem, using real data. The goal of this project is to explore the performance of multi-armed bandit algorithms to best predict the correct dosage of Warfarin for a patient *without* a trial-and-error procedure as typically employed.

Problem setting Let T be the number of time steps. At each time step t , a new patient arrives and we observe its individual feature vector $X_t \in \mathbb{R}^d$: this represents the available knowledge about the patient (e.g., gender, age, ...). The decision-maker (your algorithm) has access to K arms, where the arm represents the warfarin dosage to provide to the patient. For simplicity, we discretize the actions into $K = 3$

- Low warfarin dose: under 3mg/day
- Medium warfarin dose: 3-7 mg/day
- High warfarin dose: above 7mg/day

If the algorithm identifies the correct dosage for the patient, the reward is 0, otherwise a reward of -1 is received.

We assume that the reward for each arm depends on the patient features. For simplicity, consider a linear model where the reward for arm $i \in [K]$ for a patient with features X_t is

$$r_t(X_t, a_i) = X_t^T \beta_i + \varepsilon_{i,t}$$

where $\beta_i \in \mathbb{R}^d$ is an unknown parameter and the $\varepsilon_{i,t}$ are independent Gaussian random variables with zero mean.

The goal of this project is to design a bandit algorithm that learns a mapping $\pi : X \rightarrow a$ that yields the maximal expected reward. Let $\pi_t(s) \in [K]$ denote the arm chosen by policy π at time $t \in [T]$ for the patient s_t . Define the optimal policy π^* to be the policy that maximizes the expected reward across patients given the true β_j parameters for $j = 1, \dots, K$: this algorithm always chose the the

arm with $\max_j (X_t^T \beta_j)$. Under the linear model, if the agent chooses arm i at timestep t it will incur the expected regret of

$$\mathbb{E}[\max_j [X_t^T \beta_j] - X_t^T \beta_i]$$

The goal is to create and evaluate algorithms that minimize the cumulative expected regret $R_T = \sum_{t=1}^T \mathbb{E}[\max_j [X_t^T \beta_j] - X_t^T \beta_i]$.

Lattimore and Szepesvári have a nice series of blog posts that provide a good introduction to bandit algorithms, available here: [BanditAlgs.com](#). The [Introduction](#) and the [Linear Bandit](#) posts may be particularly of interest. For more details of the available Bandit literature you can check out the [Bandit Algorithms Book](#) by the same authors.

Dataset

We use a publicly available patient dataset that was collected by staff at the Pharmacogenetics and Pharmacogenomics Knowledge Base (PharmGKB) for 5700 patients who were treated with warfarin from 21 research groups spanning 9 countries and 4 continents. You can find the data in `warfarin.csv` and metadata containing a description of each column in `metadata.xls`. Features of each patient in this dataset includes, demographics (gender, race, ...), background (height, weight, medical history, ...), phenotypes and genotypes.

Importantly, this data contains the true patient-specific optimal warfarin doses (which are initially unknown but are eventually found through the physician-guided dose adjustment process over the course of a few weeks) for 5528 patients. You may find this data in mg/week in `Therapeutic Dose of Warfarin`¹ column in `warfarin.csv`. There are in total 5528 patient with known therapeutic dose of warfarin in the dataset (you may drop and ignore the remaining 172 patients for the purpose of this project). Given this data you can classify the right dosage for each patient as *low*: less than 21 mg/week, *medium*: 21-49 mg/week and *high*: more than 49 mg/week, as defined in [Consortium \(2009\)](#) and [Introduction](#).

Missing Data: There are missing values in the dataset, you may impute some of the missing values, or just treat "missing" or "unknown" as an extra possible feature value: for example, gender can be {male, female, unknown}. In `appx.pdf` section S4 you may find guidelines on how to impute some of the missing genotypes.

Baselines [10pts]

To become familiar with the data the first task is to measure the performance of existing approaches. As described in [Introduction](#) there exist three main approaches for determining the initial warfarin dose

1. *Fixed-dose*: This approach will assign 5mg/day (medium) dose to all patients.
2. *Warfarin Clinical Dosing Algorithm*: This method is a linear model based on age, height, weight, race and medications that patient is taking. You can find the exact model in section S1f of `appx.pdf`.

¹You cannot use `Therapeutic Dose of Warfarin` data as an input to your algorithm.

3. *Warfarin Pharmacogenetic Dosing Algorithm*: This method is also a linear model proposed by Consortium (2009), and it also includes genotypes as an input. You can find the exact model is section S1e of `appx.pdf`.

Choose two of the three models above, and report the performance of these methods on the dataset.

Performance: the performance metric is the fraction of right decisions an algorithm makes. We say an algorithm has made the right decision if the action is same as bucketed dosage in column *Therapeutic Dose of Warfarin* of the dataset.

Implementing a Linear Bandit Algorithm on Warfarin Dataset [50pts]

Your task is to implement a bandit algorithm that outperforms at least the *fixed-dose* baseline and hopefully the other one. Your bandit algorithm will start with just the knowledge of the action set and the features map (you can choose any subset of features from the dataset, you may look at section S1d in `appx.pdf` for the most relevant features. You cannot use the ground truth, *Therapeutic Dose of Warfarin*, as a feature). When a patient comes in, your bandit algorithm makes a recommendation on the amount of Warfarin for that patient (low, medium, high) and observes either a reward of 0 or -1, corresponding to a correct, respectively incorrect, decision. your algorithm will then use the observed reward to update its parameters.

To simulate an online learning environment you will need to sample (without replacement) patients from the provided dataset and pass their features to the algorithm. The regret is evaluated while the algorithm learns on the patients provided to the algorithm: you can assume that the optimal policy always makes the right decision.

Algorithm: One option to solve this question is to use a linear bandit algorithm [Linear Bandit](#), though which variant to use is up to you.

Performance: You should report two measure of performance. *Regret* as defined in section [Introduction](#), and *the average fraction of incorrect dosing decisions* under each policy as a function of the number of patients seen in the data.

Robustness: In order to make sure your performance gain/ loss over baselines is not because of ordering of the patients, run your algorithms multiple times on different random permutations of the patients and plot the performance of your algorithm with confidence intervals.

Further Directions [40pts]

There are multiple avenues you can follow, we propose some suggestions here, but you are free to chose the direction of your project from here.

- **Reward Structure:** The reward is currently defined as discrete values $\{-1, 0\}$. An interesting issue is to explore using real-valued rewards and study the performance and regret of your algorithm.
- **Risk Sensitivity:** One benefit of *fixed-dose* approach is that the distance between a patient's received dosage and their true optimal dosage has a smaller range vs alternate algorithms which may prescribe a patient that needs a low dose, a high dose. Propose to a change to your algorithm in order to make less severe mistakes, and verify if it reduces the unceratinty?

- **Other Algorithms** You can compare your algorithm to other Bandits algorithms, for example the Lasso Bandit [Bastani and Bayati \(2015\)](#) which uses a special type of regularization to improve performance. [Bastani and Bayati \(2015\)](#) used the same dataset and problem setup as you did.

References

- H. Bastani and M. Bayati. Online decision-making with high-dimensional covariates. 2015.
- I. W. P. Consortium. Estimation of the warfarin dose with clinical and pharmacogenetic data. *New England Journal of Medicine*, 360(8):753–764, 2009.